

# Biostat 203B Homework 1

Due Jan 24, 2024 @ 11:59PM

Yanzi Sun 106183069

## Table of contents

Q1. Git/GitHub . . . . .	2
Q2. Data ethics training . . . . .	2
Q3. Linux Shell Commands . . . . .	3
Q4. Who's popular in Price and Prejudice . . . . .	9
Q5. More fun with Linux . . . . .	10
Q6. Book . . . . .	11

Display machine information for reproducibility:

```
sessionInfo()
```

```
R version 4.4.2 (2024-10-31)
Platform: aarch64-apple-darwin20
Running under: macOS Sonoma 14.7.1
```

```
Matrix products: default
BLAS:   /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRblas.0.dylib
LAPACK: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRlapack.dylib;
```

```
locale:
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
```

```
time zone: America/Los_Angeles
tzcode source: internal
```

```
attached base packages:
[1] stats      graphics  grDevices  utils      datasets  methods   base
```

loaded via a namespace (and not attached):

```
[1] compiler_4.4.2    fastmap_1.2.0    cli_3.6.3        tools_4.4.2
[5] htmltools_0.5.8.1 rstudioapi_0.17.1 yaml_2.3.10      rmarkdown_2.29
[9] knitr_1.49        jsonlite_1.8.9   xfun_0.50        digest_0.6.37
[13] rlang_1.1.4       evaluate_1.0.1
```

## Q1. Git/GitHub

**No handwritten homework reports are accepted for this course.** We work with Git and GitHub. Efficient and abundant use of Git, e.g., frequent and well-documented commits, is an important criterion for grading your homework.

1. Apply for the [Student Developer Pack](#) at GitHub using your UCLA email. You'll get GitHub Pro account for free (unlimited public and private repositories).
2. Create a **private** repository `biostat-203b-2025-winter` and add Hua-Zhou and TA team (Tomoki-Okuno for Lec 1; `parsajamshidian` and `BowenZhang2001` for Lec 82) as your collaborators with write permission.
3. Top directories of the repository should be `hw1`, `hw2`, ... Maintain two branches `main` and `develop`. The `develop` branch will be your main playground, the place where you develop solution (code) to homework problems and write up report. The `main` branch will be your presentation area. Submit your homework files (Quarto file `qmd`, `html` file converted by Quarto, all code and extra data sets to reproduce results) in the `main` branch.
4. After each homework due date, course reader and instructor will check out your `main` branch for grading. Tag each of your homework submissions with tag names `hw1`, `hw2`, ... Tagging time will be used as your submission time. That means if you tag your `hw1` submission after deadline, penalty points will be deducted for late submission.
5. After this course, you can make this repository public and use it to demonstrate your skill sets on job market.

**Solution:** Done on Jan11.

## Q2. Data ethics training

This exercise (and later in this course) uses the [MIMIC-IV data v3.1](#), a freely accessible critical care database developed by the MIT Lab for Computational Physiology. Follow the instructions at <https://mimic.mit.edu/docs/gettingstarted/> to (1) complete the **CITI Data or Specimens Only Research** course and (2) obtain the PhysioNet credential for using the MIMIC-IV data. Display the verification links to your completion report and completion certificate here. **You must complete Q2 before working on the remaining questions.**

(Hint: The CITI training takes a few hours and the PhysioNet credentialing takes a couple days; do not leave it to the last minute.)

**Solution:** Here is the link of the completion report:<https://www.citiprogram.org/verify/?k1283ea47-0ca3-42b7-9d2b-e721bfb691cb-67209968>, completion certificate:<https://www.citiprogram.org/verify/?w3df54344-f850-42f3-b0df-9a632ad2eecf-67209968>

### Q3. Linux Shell Commands

1. Make the MIMIC-IV v3.1 data available at location `~/mimic`. The output of the `ls -l ~/mimic` command should be similar to the below (from my laptop).

```
# content of mimic folder
ls -l ~/mimic/
```

```
total 56
-rw-r--r--  1 yanzisun  staff   15199 Oct 10 13:29 CHANGELOG.txt
-rw-r--r--  1 yanzisun  staff    2518 Oct 10 14:30 LICENSE.txt
-rw-r--r--  1 yanzisun  staff    2884 Oct 11 14:55 SHA256SUMS.txt
drwxr-xr-x@ 24 yanzisun  staff     768 Jan 23 23:05 hosp
drwxr-xr-x@ 11 yanzisun  staff     352 Jan 21 09:25 icu
-rw-r--r--  1 yanzisun  staff     789 Jan 14 11:41 index.html
```

Refer to the documentation <https://physionet.org/content/mimiciv/3.1/> for details of data files. Do **not** put these data files into Git; they are big. Do **not** copy them into your directory. Do **not** decompress the gz data files. These create unnecessary big files and are not big-data-friendly practices. Read from the data folder `~/mimic` directly in following exercises.

Use Bash commands to answer following questions.

**Solution:** I downloaded MIMIC IV v3.1 to my computer and made it available at `~/mimic/`.

2. Display the contents in the folders `hosp` and `icu` using Bash command `ls -l`. Why are these data files distributed as `.csv.gz` files instead of `.csv` (comma separated values) files? Read the page <https://mimic.mit.edu/docs/iv/> to understand what's in each folder.

**Solution:** Below is the answer to question 2. The files are gzip-compressed csv files so they has `.gz` at the end of file name. Compressing files can save storage space.

```
ls -l ~/mimic/hosp
ls -l ~/mimic/icu
```

```

total 12306248
-rw-r--r--@ 1 yanzisun staff 19928140 Jun 24 2024 admissions.csv.gz
-rw-r--r--@ 1 yanzisun staff 427554 Apr 12 2024 d_hcpcs.csv.gz
-rw-r--r--@ 1 yanzisun staff 876360 Apr 12 2024 d_icd_diagnoses.csv.gz
-rw-r--r--@ 1 yanzisun staff 589186 Apr 12 2024 d_icd_procedures.csv.gz
-rw-r--r--@ 1 yanzisun staff 13169 Oct 3 09:07 d_labitems.csv.gz
-rw-r--r--@ 1 yanzisun staff 33564802 Oct 3 09:07 diagnoses_icd.csv.gz
-rw-r--r--@ 1 yanzisun staff 9743908 Oct 3 09:07 drgcodes.csv.gz
-rw-r--r--@ 1 yanzisun staff 811305629 Apr 12 2024 emar.csv.gz
-rw-r--r--@ 1 yanzisun staff 748158322 Apr 12 2024 emar_detail.csv.gz
-rw-r--r--@ 1 yanzisun staff 2162335 Apr 12 2024 hcpcsevents.csv.gz
-rw-r--r--@ 1 yanzisun staff 2592909134 Oct 3 09:08 labevents.csv.gz
-rw-r--r--@ 1 yanzisun staff 117644075 Oct 3 09:08 microbiologyevents.csv.gz
-rw-r--r--@ 1 yanzisun staff 44069351 Oct 3 09:08 omr.csv.gz
-rw-r--r--@ 1 yanzisun staff 2835586 Apr 12 2024 patients.csv.gz
-rw-r--r--@ 1 yanzisun staff 525708076 Apr 12 2024 pharmacy.csv.gz
-rw-r--r--@ 1 yanzisun staff 666594177 Apr 12 2024 poe.csv.gz
-rw-r--r--@ 1 yanzisun staff 55267894 Apr 12 2024 poe_detail.csv.gz
-rw-r--r--@ 1 yanzisun staff 606298611 Apr 12 2024 prescriptions.csv.gz
-rw-r--r--@ 1 yanzisun staff 7777324 Apr 12 2024 procedures_icd.csv.gz
-rw-r--r--@ 1 yanzisun staff 127330 Apr 12 2024 provider.csv.gz
-rw-r--r--@ 1 yanzisun staff 8569241 Apr 12 2024 services.csv.gz
-rw-r--r--@ 1 yanzisun staff 46185771 Oct 3 09:08 transfers.csv.gz
total 8506784
-rw-r--r--@ 1 yanzisun staff 41566 Apr 12 2024 caregiver.csv.gz
-rw-r--r--@ 1 yanzisun staff 3502392765 Apr 12 2024 chartevents.csv.gz
-rw-r--r--@ 1 yanzisun staff 58741 Apr 12 2024 d_items.csv.gz
-rw-r--r--@ 1 yanzisun staff 63481196 Apr 12 2024 datatimeevents.csv.gz
-rw-r--r--@ 1 yanzisun staff 3342355 Oct 3 07:36 icustays.csv.gz
-rw-r--r--@ 1 yanzisun staff 311642048 Apr 12 2024 ingredientevents.csv.gz
-rw-r--r--@ 1 yanzisun staff 401088206 Apr 12 2024 inputevents.csv.gz
-rw-r--r--@ 1 yanzisun staff 49307639 Apr 12 2024 outputevents.csv.gz
-rw-r--r--@ 1 yanzisun staff 24096834 Apr 12 2024 procedureevents.csv.gz

```

3. Briefly describe what Bash commands `zcat`, `zless`, `zmore`, and `zgrep` do.

**Solution:** `zcat` allows to view a compressed (zipped) file directly, basically a `cat` command for zipped files. `zless` view compressed file contents in a paginated way, can also do search option; `zmore` view compressed file contents in a new tab; `zgrep` search string/pattern in the compressed file.

4. (Looping in Bash) What's the output of the following bash script?

```
for datafile in ~/mimic/hosp/{a,l,pa}*.gz
do
    ls -l $datafile
done
```

**Solution:** The results are a long list(including file permission and sizes) of all files starting with “a,l, or pa” inside the hosp folder.

Display the number of lines in each data file using a similar loop. (Hint: combine linux commands `zcat` < and `wc -l`.)

**Solution:**

```
for datafile in ~/mimic/hosp/{a,l,pa}*.gz
do
    zcat < $datafile | wc -l
done
```

```
546029
158374765
364628
```

5. Display the first few lines of `admissions.csv.gz`. How many rows are in this data file, excluding the header line? Each `hadm_id` identifies a hospitalization. How many hospitalizations are in this data file? How many unique patients (identified by `subject_id`) are in this data file? Do they match the number of patients listed in the `patients.csv.gz` file? (Hint: combine Linux commands `zcat` <, `head/tail`, `awk`, `sort`, `uniq`, `wc`, and `so on`.)

**Solution:** There are 546028 rows in this file excluding the headerline. There are 546028 hospitalizations in this data file. There are 223452 unique patients, which does not match the number of patients listed in the `patients.csv`(364627)

```
zcat < ~/mimic/hosp/admissions.csv.gz | head -5
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | wc -l
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F',' '{print $2}' | wc -l
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F',' '{print $1}' | uniq | wc -l
zcat < ~/mimic/hosp/patients.csv.gz | tail -n +2 | uniq | wc -l
```

```
subject_id,hadm_id,admittime,dischtime,deathtime,admission_type,admit_provider_id,admission_
10000032,22595853,2180-05-06 22:23:00,2180-05-07 17:15:00,,URGENT,P49AFC,TRANSFER FROM HOSPI
10000032,22841357,2180-06-26 18:27:00,2180-06-27 18:49:00,,EW EMER.,P784FA,EMERGENCY ROOM,HOI
```

```

10000032,25742920,2180-08-05 23:44:00,2180-08-07 17:50:00,,EW EMER.,P19UTS,EMERGENCY ROOM,HOS
10000032,29079034,2180-07-23 12:35:00,2180-07-25 17:55:00,,EW EMER.,P060TX,EMERGENCY ROOM,HOS
546028
546028
223452
364627

```

6. What are the possible values taken by each of the variable `admission_type`, `admission_location`, `insurance`, and `ethnicity`? Also report the count for each unique value of these variables in decreasing order. (Hint: combine Linux commands `zcat`, `head/tail`, `awk`, `uniq -c`, `wc`, `sort`, and so on; skip the header line.)

**Solution:** Column 6,8,10,13 corresponds to `admission_type`, `admission_location`, `insurance`, and `ethnicity`.

```

for col in 6 8 10 13;
do
echo "Count and sorted unique values of variable $col:"
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F, -v c=$col '{print $c}' | sort
echo
done

```

Count and sorted unique values of variable 6:

```

177459 EW EMER.
119456 EU OBSERVATION
84437 OBSERVATION ADMIT
54929 URGENT
42898 SURGICAL SAME DAY ADMISSION
24551 DIRECT OBSERVATION
21973 DIRECT EMER.
13130 ELECTIVE
7195 AMBULATORY OBSERVATION

```

Count and sorted unique values of variable 8:

```

244179 EMERGENCY ROOM
163228 PHYSICIAN REFERRAL
56227 TRANSFER FROM HOSPITAL
42365 WALK-IN/SELF REFERRAL
12965 CLINIC REFERRAL
8518 PROCEDURE SITE
6317 TRANSFER FROM SKILLED NURSING FACILITY
5837 INTERNAL TRANSFER TO OR FROM PSYCH
5734 PACU

```

402 INFORMATION NOT AVAILABLE  
255 AMBULATORY SURGERY TRANSFER  
1

Count and sorted unique values of variable 10:

244576 Medicare  
173399 Private  
104229 Medicaid  
14006 Other  
9355  
463 No charge

Count and sorted unique values of variable 13:

336538 WHITE  
75482 BLACK/AFRICAN AMERICAN  
19788 OTHER  
13972 WHITE - OTHER EUROPEAN  
13870 UNKNOWN  
10903 HISPANIC/LATINO - PUERTO RICAN  
8287 HISPANIC OR LATINO  
7809 ASIAN  
7644 ASIAN - CHINESE  
6597 WHITE - RUSSIAN  
6205 BLACK/CAPE VERDEAN  
6070 HISPANIC/LATINO - DOMINICAN  
3875 BLACK/CARIBBEAN ISLAND  
3495 BLACK/AFRICAN  
3478 UNABLE TO OBTAIN  
2162 PATIENT DECLINED TO ANSWER  
2082 PORTUGUESE  
1973 ASIAN - SOUTH EAST ASIAN  
1886 WHITE - EASTERN EUROPEAN  
1858 HISPANIC/LATINO - GUATEMALAN  
1661 ASIAN - ASIAN INDIAN  
1526 WHITE - BRAZILIAN  
1320 HISPANIC/LATINO - SALVADORAN  
1247 AMERICAN INDIAN/ALASKA NATIVE  
920 HISPANIC/LATINO - COLUMBIAN  
883 HISPANIC/LATINO - MEXICAN  
774 SOUTH AMERICAN  
725 HISPANIC/LATINO - HONDURAN  
664 ASIAN - KOREAN  
641 HISPANIC/LATINO - CUBAN

603 HISPANIC/LATINO - CENTRAL AMERICAN  
596 MULTIPLE RACE/ETHNICITY  
494 NATIVE HAWAIIAN OR OTHER PACIFIC ISLANDER

7. The `icustays.csv.gz` file contains all the ICU stays during the study period. How many ICU stays, identified by `stay_id`, are in this data file? How many unique patients, identified by `subject_id`, are in this data file?

**Solution:** There are 94459 stays and 65367 unique patients in this data file.

```
zcat < ~/mimic/icu/icustays.csv.gz | awk -F ',' '{print $3}' | wc -l  
zcat < ~/mimic/icu/icustays.csv.gz | awk -F ',' '{print $1}' | uniq | wc -l
```

94459  
65367

8. *To compress, or not to compress. That's the question.* Let's focus on the big data file `labevents.csv.gz`. Compare compressed gz file size to the uncompressed file size. Compare the run times of `zcat < ~/mimic/labevents.csv.gz | wc -l` versus `wc -l labevents.csv`. Discuss the trade off between storage and speed for big data files. (Hint: `gzip -dk < FILENAME.gz > ./FILENAME`. Remember to delete the large `labevents.csv` file after the exercise.)

**Solution:** From the results I get, I see the runtimes are basically the same. however, the unzipped .csv file takes much larger storage than its zipped .csv.gz form. I will prefer to use `zcat` for compressed files.

```
time zcat < ~/mimic/hosp/labevents.csv.gz | wc -l  
gzip -k -d ~/mimic/hosp/labevents.csv.gz  
time wc -l ~/mimic/hosp/labevents.csv  
rm ~/mimic/hosp/labevents.csv
```

158374765

```
real    0m18.724s  
user    0m29.245s  
sys 0m1.795s  
158374765 /Users/yanzisun/mimic/hosp/labevents.csv
```

```
real    0m19.497s  
user    0m17.492s  
sys 0m1.741s
```



#### Q4. Who's popular in Price and Prejudice

1. You and your friend just have finished reading *Pride and Prejudice* by Jane Austen. Among the four main characters in the book, Elizabeth, Jane, Lydia, and Darcy, your friend thinks that Darcy was the most mentioned. You, however, are certain it was Elizabeth. Obtain the full text of the novel from <http://www.gutenberg.org/cache/epub/42671/pg42671.txt> and save to your local folder.

```
wget -nc http://www.gutenberg.org/cache/epub/42671/pg42671.txt
```

Explain what `wget -nc` does. Do **not** put this text file `pg42671.txt` in Git. Complete the following loop to tabulate the number of times each of the four characters is mentioned using Linux commands.

**Solution:** ‘`wget -nc`’ downloads file from websites, `-nc` is the option that specifies not to create duplicate if the file was downloaded already. The counts for each names are calculated below.

```
for char in Elizabeth Jane Lydia Darcy
do
    echo $char:
    grep -c "$char" pg42671.txt
done
```

```
Elizabeth:
633
Jane:
289
Lydia:
166
Darcy:
414
```

2. What's the difference between the following two commands?

```
echo 'hello, world' > test1.txt
```

and

```
echo 'hello, world' >> test2.txt
```

**Solution:** `>` redirects the standard output to a file. Specifically, `>` save text to output file (overwrite if file exist); `>>` append to the end of the output file and saves it.

3. Using your favorite text editor (e.g., vi), type the following and save the file as `middle.sh`:

```
#!/bin/sh
# Select lines from the middle of a file.
# Usage: bash middle.sh filename end_line num_lines
head -n "$2" "$1" | tail -n "$3"
```

Using `chmod` to make the file executable by the owner, and run

```
./middle.sh pg42671.txt 20 5
```

Release date: May 9, 2013 [eBook #42671]

Language: English

Explain the output. Explain the meaning of "\$1", "\$2", and "\$3" in this shell script. Why do we need the first line of the shell script?

**Solution:** I used this `chmod` command line to edit user access: `chmod u+x middle.sh`. The output is the result of running shell script written in `[middle.sh]` on `[pg42671.txt]`. The "\$1" represents the filename, "\$2" represents the `end_line`, and "\$3" represents the number of lines. Basically the `head` command is translated to: `head -n endline filename | tail -n num lines`. When the shell script is executable and applied to the `txt` file, it extracts line 15 to 20 as the end line is 20 and the number of lines are 5.

## Q5. More fun with Linux

Try following commands in Bash and interpret the results: `cal`, `cal 2025`, `cal 9 1752` (anything unusual?), `date`, `hostname`, `arch`, `uname -a`, `uptime`, `who am i`, `who`, `w`, `id`, `last | head`, `echo {con,pre}{sent,fer}{s,ed}`, `time sleep 5`, `history | tail`.

**Solution:** I played with all commands and # them so they do not run multiple times when I test the other commands. `cal` gives the calendar of the year or month indicated. `date` gives the current date and time based off your system time zone. `hostname` gives the name of my PC. `arch` gives my PC's architecture type. `uname -a` gives detail information of my PC, including kernel, hostname, operating system, time, root, and architecture. `uptime` gives current system time, how long the PC has been running since last restart, average load time, and number of users. `who am i` tells the user is me and gives system time. `who` gives a list of users and time using the PC. `w` provides a detailed use history of past users. `id` gives information about the user identity. `last` provides a list of logins into the PC, `head` option limits the output number

to 10. `echo {con,pre}{sent,fer}{s,ed}` prints out all possible combinations of the characters in curly brackets. `2x2x2=8` the sleep command pauses execution for 5 seconds and the time command measures how long it actually takes. `history | tail` displays the last 10 commands from my shell history.

```
#cal 2025
#cal 1 2001
#date
#hostname
arch
#uname -a
#uptime
#who am i
who
w
#id
#last |head
#echo {con,pre}{sent,fer}{s,ed}
#time sleep 5
#history | tail
```

```
arm64yanzisun      ttys000      Jan 23 15:34
yanzisun           ttys001      Jan 23 15:34
yanzisun           ttys002      Jan 23 15:34
yanzisun           ttys003      Jan 23 17:31
yanzisun           console      Jan 23 15:34
23:09 up 20 days, 1:26, 5 users, load averages: 2.03 2.08 2.07
USER      TTY      FROM      LOGIN@    IDLE WHAT
yanzisun  s000      -          15:34     7:35 -bash
yanzisun  s001      -          15:34     7:35 -bash
yanzisun  s002      -          15:34     6:01 -bash
yanzisun  s003      -          17:31     7 -bash
yanzisun  console  -          15:34     7:35 -
```

## Q6. Book

1. Git clone the repository <https://github.com/christophergandrud/Rep-Res-Book> for the book *Reproducible Research with R and RStudio* to your local machine. Do **not** put this repository within your homework repository `biostat-203b-2025-winter`.

**Solution:** Done. I # them because rendering it on my local creates duplicates everytime running this cell. If you want to run it, please remove #.

```
#git init
#git clone https://github.com/christophergandrud/Rep-Res-Book
```

2. Open the project by clicking `rep-res-3rd-edition.Rproj` and compile the book by clicking **Build Book** in the **Build** panel of RStudio. (Hint: I was able to build `git_book` and `epub_book` directly. For `pdf_book`, I needed to add a line `\usepackage{hyperref}` to the file `Rep-Res-Book/rep-res-3rd-edition/latex/preabmle.tex`.)

The point of this exercise is (1) to obtain the book for free and (2) to see an example how a complicated project such as a book can be organized in a reproducible way. Use `sudo apt install PKGNAME` to install required Ubuntu packages and `tlmgr install PKGNAME` to install missing TeXLive packages.

For grading purpose, include a screenshot of Section 4.1.5 of the book here.

**Solution:** I cloned the repository and compiled the book. here is the screenshot of 4.1.5 of the

