

Biostat 203B Homework 5

Due Mar 20 @ 11:59PM

Yanzi Sun 106183069

Table of contents

Predicting ICU duration	1
-----------------------------------	---

Predicting ICU duration

Using the ICU cohort `mimiciv_icu_cohort.rds` you built in Homework 4, develop at least three machine learning approaches (logistic regression with enet regularization, random forest, boosting, SVM, MLP, etc) plus a model stacking approach for predicting whether a patient's ICU stay will be longer than 2 days. You should use the `los_long` variable as the outcome. Your algorithms can use patient demographic information (gender, age at ICU `intime`, marital status, race), ICU admission information (first care unit), the last lab measurements before the ICU stay, and first vital measurements during ICU stay as features. You are welcome to use any feature engineering techniques you think are appropriate; but make sure to not use features that are not available at an ICU stay's `intime`. For instance, `last_careunit` cannot be used in your algorithms.

1. Data preprocessing and feature engineering.

```
#load necessary libraries
library(GGally)
```

Loading required package: `ggplot2`

```
Registered S3 method overwritten by 'GGally':
  method from
+.gg      ggplot2
```

```
library(gtsummary)
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v lubridate  1.9.4      v tibble     3.2.1
v purrr      1.0.2      v tidyr      1.3.1
```

```
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
library(tidymodels)
```

```
-- Attaching packages ----- tidymodels 1.2.0 --
v broom       1.0.7      v rsample     1.2.1
v dials       1.3.0      v tune        1.2.1
v infer       1.0.7      v workflows   1.1.4
v modeldata   1.4.0      v workflowsets 1.1.0
v parsnip     1.2.1      v yardstick   1.3.2
v recipes     1.1.0
```

```
-- Conflicts ----- tidymodels_conflicts() --
x scales::discard() masks purrr::discard()
x dplyr::filter()   masks stats::filter()
x recipes::fixed()  masks stringr::fixed()
x dplyr::lag()       masks stats::lag()
x yardstick::spec() masks readr::spec()
x recipes::step()    masks stats::step()
* Dig deeper into tidy modeling with R at https://www.tmw.r.org
```

```
library(tensorflow)
library(keras)
```

Attaching package: 'keras'

The following object is masked from 'package:yardstick':

get_weights

```
library(reticulate)
library(xgboost)
```

Attaching package: 'xgboost'

The following object is masked from 'package:dplyr':

slice

```
library(ranger)
#load the mimic-icu dataset
mimic_icu_cohort <- readRDS("~/Desktop/203B/203B-HW/hw4/mimiciv_shiny/mimic_icu_cohort.rds")
mutate(
  insurance = as.factor(insurance),
  marital_status = as.factor(marital_status),
  gender = as.factor(gender),
  los_long = as.logical(los_long), # Ensure it is TRUE/FALSE
  los_long = factor(los_long, levels = c(FALSE, TRUE),
                    labels = c("no", "yes"))) |>
drop_na(los_long) #|>
#slice_sample(prop = 0.25) # downsample for faster computation

#have an overview of the dataset
#mimic_icu_cohort |> tbl_summary(by = los_long)
```

2. Partition data into 50% training set and 50% test set. Stratify partitioning according to `los_long`. For grading purpose, sort the data by `subject_id`, `hadm_id`, and `stay_id` and use the seed 203 for the initial data split. Below is the sample code.

```
##| eval: false
set.seed(203)

# sort
mimic_icu_cohort <- mimic_icu_cohort |>
  arrange(subject_id, hadm_id, stay_id) |>
  select(first_careunit, gender, age_at_intime, marital_status, race,
         `heart rate`, `non invasive blood pressure systolic`,
         `non invasive blood pressure diastolic`,
         `respiratory rate`, `temperature fahrenheit`,
         bicarbonate, chloride, creatinine, glucose, potassium,
```

```

        sodium, hematocrit, wbc,
        los_long)

#Initial split into test and non-test sets

data_split <- initial_split(
  mimic_icu_cohort,
  # stratify by los_long
  strata = "los_long",
  prop = 0.5
)

mimic_other <- training(data_split)
dim(mimic_other)

```

```
[1] 47221    19
```

```

mimic_test <- testing(data_split)
dim(mimic_test)

```

```
[1] 47223    19
```

3. Train and tune the models using the training set.

Solution: I choose to do logistic regression, SVM, xgboost, and finally model stacking.

```

#this chunk is checking the data cleanliness, do not need to run so I commented
# mimic_other %>%
#   summarise(across(everything(), ~ sum(is.na(.)))) %>%
#   print(width = Inf)
#
# table(mimic_icu_cohort$los_long)
# summary(mimic_icu_cohort$los_long)
# str(mimic_icu_cohort$los_long)
# table(mimic_other$los_long)

```

##Here is the code for logistic regression:

```

#Recipe
logit_recipe <-
  recipe(
    los_long ~ .,
    data = mimic_other
  ) |>
  step_unknown(all_nominal(), -all_outcomes()) |> # Handle missing categorical values
  # mean imputation for numeric variable
  step_impute_mean(`heart rate`, `non invasive blood pressure systolic`, `non invasive blood
    `respiratory rate`, `temperature fahrenheit`,
    bicarbonate, chloride, creatinine, glucose, potassium, sodium,
    hematocrit, wbc) |>
  # mode imputation for categorical variable
  step_impute_mode(marital_status) |>
  # create traditional dummy variables
  step_dummy(all_nominal_predictors()) |>
  # zero-variance filter
  step_zv(all_numeric_predictors()) |>
  # center and scale numeric data
  step_normalize(all_numeric_predictors()) |>
  # estimate the means and standard deviations
  # prep(training = Heart_other, retain = TRUE) |>
  print()

```

-- Recipe -----

-- Inputs

Number of variables by role

```

outcome:    1
predictor: 18

```

-- Operations

```

* Unknown factor level assignment for: all_nominal() -all_outcomes()

* Mean imputation for: `heart rate`, ...

* Mode imputation for: marital_status

* Dummy variables from: all_nominal_predictors()

* Zero variance filter on: all_numeric_predictors()

* Centering and scaling for: all_numeric_predictors()

```

```

# Model
logit_mod <-
  logistic_reg(
    penalty = tune(),
    mixture = tune()
  ) |>
  set_engine("glmnet", standardize = FALSE) |>
  print()

```

Logistic Regression Model Specification (classification)

Main Arguments:

```

  penalty = tune()
  mixture = tune()

```

Engine-Specific Arguments:

```

  standardize = FALSE

```

Computational engine: glmnet

```

# Workflow, bundle the recipe and model

logit_wf <- workflow() |>
  add_recipe(logit_recipe) |>
  add_model(logit_mod) |>
  print()

```

```

== Workflow =====
Preprocessor: Recipe
Model: logistic_reg()

-- Preprocessor -----
6 Recipe Steps

* step_unknown()
* step_impute_mean()
* step_impute_mode()
* step_dummy()
* step_zv()
* step_normalize()

-- Model -----
Logistic Regression Model Specification (classification)

Main Arguments:
  penalty = tune()
  mixture = tune()

Engine-Specific Arguments:
  standardize = FALSE

Computational engine: glmnet

# Tune the model(the penalty and mixture hyperparameters).

param_grid <- grid_regular(
  penalty(range = c(-6, 3)),
  mixture(),
  levels = c(100, 5)
) |>
print()

# A tibble: 500 x 2
  penalty mixture
    <dbl>    <dbl>
1 0.000001      0
2 0.00000123    0
3 0.00000152    0
4 0.00000187    0

```

```

5 0.00000231      0
6 0.00000285      0
7 0.00000351      0
8 0.00000433      0
9 0.00000534      0
10 0.00000658     0
# i 490 more rows

```

```

# Set Cross-Validation partitions
set.seed(203)

folds <- vfold_cv(mimic_other, v = 5)
folds

```

```

# 5-fold cross-validation
# A tibble: 5 x 2
  splits      id
  <list>     <chr>
1 <split [37776/9445]> Fold1
2 <split [37777/9444]> Fold2
3 <split [37777/9444]> Fold3
4 <split [37777/9444]> Fold4
5 <split [37777/9444]> Fold5

```

```

# fit cross-validation
(logit_fit <- logit_wf |>
  tune_grid(
    resamples = folds,
    grid = param_grid,
    metrics = metric_set(roc_auc, accuracy)
  ) |>
  system.time()

```

```

      user system elapsed
29.180   2.178  31.417

```

```
logit_fit
```

```

# Tuning results
# 5-fold cross-validation

```



```
# A tibble: 5 x 4
  splits          id    .metrics          .notes
  <list>         <chr> <list>          <list>
1 <split [37776/9445]> Fold1 <tibble [1,000 x 6]> <tibble [0 x 3]>
2 <split [37777/9444]> Fold2 <tibble [1,000 x 6]> <tibble [0 x 3]>
3 <split [37777/9444]> Fold3 <tibble [1,000 x 6]> <tibble [0 x 3]>
4 <split [37777/9444]> Fold4 <tibble [1,000 x 6]> <tibble [0 x 3]>
5 <split [37777/9444]> Fold5 <tibble [1,000 x 6]> <tibble [0 x 3]>
```


```
# Visualize
logit_fit |>
  # aggregate metrics from K folds
  collect_metrics() |>
  print(width = Inf) |>
  filter(.metric == "roc_auc") |>
  ggplot(mapping = aes(x = penalty, y = mean, color = factor(mixture))) +
  geom_point() +
  labs(x = "Penalty", y = "CV AUC") +
  scale_x_log10()
```

```
# A tibble: 1,000 x 8
  penalty mixture .metric .estimator mean    n std_err
  <dbl>   <dbl> <chr>   <chr>    <dbl> <int>  <dbl>
1 0.000001      0 accuracy binary    0.562     5 0.00230
2 0.000001      0 roc_auc  binary    0.588     5 0.00246
3 0.00000123    0 accuracy binary    0.562     5 0.00230
4 0.00000123    0 roc_auc  binary    0.588     5 0.00246
5 0.00000152    0 accuracy binary    0.562     5 0.00230
6 0.00000152    0 roc_auc  binary    0.588     5 0.00246
7 0.00000187    0 accuracy binary    0.562     5 0.00230
8 0.00000187    0 roc_auc  binary    0.588     5 0.00246
9 0.00000231    0 accuracy binary    0.562     5 0.00230
10 0.00000231    0 roc_auc  binary    0.588     5 0.00246
  .config
  <chr>
1 Preprocessor1_Model001
2 Preprocessor1_Model001
3 Preprocessor1_Model002
4 Preprocessor1_Model002
5 Preprocessor1_Model003
6 Preprocessor1_Model003
7 Preprocessor1_Model004
```

```

8 Preprocessor1_Model004
9 Preprocessor1_Model005
10 Preprocessor1_Model005
# i 990 more rows

```



```

#show top 5 models and select the best one
logit_fit |>
  show_best(metric = "roc_auc")

```

```
# A tibble: 5 x 8
```

	penalty	mixture	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	0.00811	0	roc_auc	binary	0.588	5	0.00246	Preprocessor1_Model~
2	0.000001	0	roc_auc	binary	0.588	5	0.00246	Preprocessor1_Model~
3	0.00000123	0	roc_auc	binary	0.588	5	0.00246	Preprocessor1_Model~
4	0.00000152	0	roc_auc	binary	0.588	5	0.00246	Preprocessor1_Model~
5	0.00000187	0	roc_auc	binary	0.588	5	0.00246	Preprocessor1_Model~

```

best_logit <- logit_fit |>
  select_best(metric = "roc_auc")
best_logit

```

```
# A tibble: 1 x 3
  penalty mixture .config
  <dbl>   <dbl> <chr>
1 0.00811      0 Preprocessor1_Model044
```

```
#finalize workflow
# Final workflow
final_wf <- logit_wf |>
  finalize_workflow(best_logit)
final_wf
```

```
== Workflow =====
Preprocessor: Recipe
Model: logistic_reg()
```

```
-- Preprocessor -----
6 Recipe Steps
```

```
* step_unknown()
* step_impute_mean()
* step_impute_mode()
* step_dummy()
* step_zv()
* step_normalize()
```

```
-- Model -----
Logistic Regression Model Specification (classification)
```

```
Main Arguments:
  penalty = 0.00811130830789687
  mixture = 0
```

```
Engine-Specific Arguments:
  standardize = FALSE
```

```
Computational engine: glmnet
```

```
# Fit the whole training set, then predict the test cases
final_fit <-
  final_wf |>
  last_fit(data_split)
final_fit
```

```
# Resampling results
# Manual resampling
# A tibble: 1 x 6
  splits          id      .metrics .notes  .predictions .workflow
  <list>         <chr>    <list>  <list>  <list>       <list>
1 <split [47221/47223]> train/test sp~ <tibble> <tibble> <tibble>   <workflow>
```

```
# Test metrics
final_fit |>
  collect_metrics()
```

```
# A tibble: 3 x 4
  .metric      .estimator .estimate .config
  <chr>       <chr>      <dbl> <chr>
1 accuracy    binary      0.561 Preprocessor1_Model1
2 roc_auc     binary      0.585 Preprocessor1_Model1
3 brier_class binary      0.245 Preprocessor1_Model1
```

##Here is the code for SVM:

```
# Recipe
library(kernlab)
```

Attaching package: 'kernlab'

The following object is masked from 'package:scales':

alpha

The following object is masked from 'package:purrr':

cross

The following object is masked from 'package:ggplot2':

alpha

```

svm_recipe <-
  recipe(
    los_long ~ .,
    data = mimic_other
  ) |>
  step_unknown(all_nominal(), -all_outcomes()) |> # Handle missing categorical values
  # mean imputation for numeric variable
  step_impute_mean(`heart rate`, `non invasive blood pressure systolic`, `non invasive blood
    `respiratory rate`, `temperature fahrenheit`,
    bicarbonate, chloride, creatinine, glucose, potassium, sodium,
    hematocrit, wbc) |>
  # mode imputation for categorical variable
  step_impute_mode(marital_status) |>
  # create traditional dummy variables
  step_dummy(all_nominal_predictors()) |>
  # zero-variance filter
  step_zv(all_numeric_predictors()) |>
  # center and scale numeric data
  step_normalize(all_numeric_predictors()) |>
  # estimate the means and standard deviations
  # prep(training = Heart_other, retain = TRUE) |>
  print()

```

-- Recipe -----

-- Inputs

Number of variables by role

```

outcome:    1
predictor: 18

```

-- Operations

- * Unknown factor level assignment for: all_nominal() -all_outcomes()
- * Mean imputation for: `heart rate`, ...
- * Mode imputation for: marital_status
- * Dummy variables from: all_nominal_predictors()
- * Zero variance filter on: all_numeric_predictors()
- * Centering and scaling for: all_numeric_predictors()

```
# Model
svm_mod <-
  svm_rbf(
    mode = "classification",
    cost = tune(),
    rbf_sigma = tune()
  ) |>
  set_engine("kernlab")
svm_mod
```

Radial Basis Function Support Vector Machine Model Specification (classification)

Main Arguments:

```
cost = tune()
rbf_sigma = tune()
```

Computational engine: kernlab

```
# Bundle recipe & Model
svm_wf <- workflow() |>
  add_recipe(svm_recipe) |>
  add_model(svm_mod)
svm_wf
```

```
== Workflow =====
Preprocessor: Recipe
Model: svm_rbf()
```

```
-- Preprocessor -----  
6 Recipe Steps
```

```
* step_unknown()  
* step_impute_mean()  
* step_impute_mode()  
* step_dummy()  
* step_zv()  
* step_normalize()
```

```
-- Model -----  
Radial Basis Function Support Vector Machine Model Specification (classification)
```

Main Arguments:

```
cost = tune()  
rbf_sigma = tune()
```

Computational engine: kernlab

```
# Tune the model  
# Trail #1 with best roc-auc ~0.593, since the best point is on the top-right edge, tune the  
# For Trail #2, reduced the sample size to 1/4 of original size to speed up the computation  
param_grid2 <- grid_regular(  
  cost(range = c(0, 6)),  
  rbf_sigma(range = c(-4, -1)),  
  levels = c(3, 3)  
)  
# param_grid1 <- grid_regular(  
#   cost(range = c(-8, 5)),  
#   rbf_sigma(range = c(-5, -3)),  
#   levels = c(3, 3)  
# )  
param_grid2
```

```
# A tibble: 9 x 2  
  cost rbf_sigma  
  <dbl>   <dbl>  
1     1  0.0001  
2     8  0.0001  
3    64  0.0001  
4     1  0.00316
```

```

5      8    0.00316
6     64    0.00316
7      1     0.1
8      8     0.1
9     64     0.1

```

```

# CV, since SVM is slow, I choose to use only 2 folds here.
set.seed(203)

folds <- vfold_cv(mimic_other, v = 2)
folds

```

```

# 2-fold cross-validation
# A tibble: 2 x 2
  splits          id
  <list>         <chr>
1 <split [23610/23611]> Fold1
2 <split [23611/23610]> Fold2

```

```

# Fit cross-validation
svm_fit <- svm_wf |>
  tune_grid(
    resamples = folds,
    grid = param_grid2,
    metrics = metric_set(roc_auc, accuracy)
  )
svm_fit

```

```

# Tuning results
# 2-fold cross-validation
# A tibble: 2 x 4
  splits          id    .metrics      .notes
  <list>         <chr> <list>      <list>
1 <split [23610/23611]> Fold1 <tibble [18 x 6]> <tibble [0 x 3]>
2 <split [23611/23610]> Fold2 <tibble [18 x 6]> <tibble [0 x 3]>

```

```

# Visualize CV results
svm_fit |>
  collect_metrics() |>
  print(width = Inf) |>
  filter(.metric == "roc_auc") |>


```



```
ggplot(mapping = aes(x = cost, y = mean, color = factor(rbf_sigma))) +
  geom_point() +
  labs(x = "Cost", y = "CV AUC") +
  scale_x_log10()
```

A tibble: 18 x 8

	cost	rbf_sigma	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	1	0.0001	accuracy	binary	0.548	2	0.00154	Preprocessor1_Model11
2	1	0.0001	roc_auc	binary	0.585	2	0.00620	Preprocessor1_Model11
3	8	0.0001	accuracy	binary	0.559	2	0.00340	Preprocessor1_Model12
4	8	0.0001	roc_auc	binary	0.589	2	0.00551	Preprocessor1_Model12
5	64	0.0001	accuracy	binary	0.563	2	0.00414	Preprocessor1_Model13
6	64	0.0001	roc_auc	binary	0.591	2	0.00591	Preprocessor1_Model13
7	1	0.00316	accuracy	binary	0.569	2	0.00372	Preprocessor1_Model14
8	1	0.00316	roc_auc	binary	0.600	2	0.00623	Preprocessor1_Model14
9	8	0.00316	accuracy	binary	0.575	2	0.00484	Preprocessor1_Model15
10	8	0.00316	roc_auc	binary	0.606	2	0.00643	Preprocessor1_Model15
11	64	0.00316	accuracy	binary	0.575	2	0.00376	Preprocessor1_Model16
12	64	0.00316	roc_auc	binary	0.606	2	0.00519	Preprocessor1_Model16
13	1	0.1	accuracy	binary	0.568	2	0.00219	Preprocessor1_Model17
14	1	0.1	roc_auc	binary	0.596	2	0.00387	Preprocessor1_Model17
15	8	0.1	accuracy	binary	0.553	2	0.00224	Preprocessor1_Model18
16	8	0.1	roc_auc	binary	0.571	2	0.00128	Preprocessor1_Model18
17	64	0.1	accuracy	binary	0.536	2	0.000837	Preprocessor1_Model19
18	64	0.1	roc_auc	binary	0.550	2	0.000656	Preprocessor1_Model19



```
svm_fit |>
  show_best(metric = "roc_auc")
```

A tibble: 5 x 8

	cost	rbf_sigma	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	64	0.00316	roc_auc	binary	0.606	2	0.00519	Preprocessor1_Model6
2	8	0.00316	roc_auc	binary	0.606	2	0.00643	Preprocessor1_Model5
3	1	0.00316	roc_auc	binary	0.600	2	0.00623	Preprocessor1_Model4
4	1	0.1	roc_auc	binary	0.596	2	0.00387	Preprocessor1_Model7
5	64	0.0001	roc_auc	binary	0.591	2	0.00591	Preprocessor1_Model3

```
best_svm <- svm_fit |>
  select_best(metric = "roc_auc")
best_svm
```

A tibble: 1 x 3

	cost	rbf_sigma	.config
	<dbl>	<dbl>	<chr>
1	64	0.00316	Preprocessor1_Model6

```
# Final workflow
final_wf <- svm_wf |>
  finalize_workflow(best_svm)
final_wf
```

```
== Workflow =====
Preprocessor: Recipe
Model: svm_rbf()

-- Preprocessor -----
6 Recipe Steps

* step_unknown()
* step_impute_mean()
* step_impute_mode()
* step_dummy()
* step_zv()
* step_normalize()

-- Model -----
Radial Basis Function Support Vector Machine Model Specification (classification)

Main Arguments:
  cost = 64
  rbf_sigma = 0.00316227766016838

Computational engine: kernlab
```

```
# Fit the whole training set, then predict the test cases
final_fit <-
  final_wf |>
  last_fit(data_split)
final_fit
```

```
# Resampling results
# Manual resampling
# A tibble: 1 x 6
```

	splits	id	.metrics	.notes	.predictions	.workflow
	<list>	<chr>	<list>	<list>	<list>	<list>
1	<split [47221/47223]>	train/test sp~	<tibble>	<tibble>	<tibble>	<workflow>

```
# Test metrics
final_fit |>
  collect_metrics()
```

```
# A tibble: 3 x 4
  .metric      .estimator .estimate .config
  <chr>        <chr>         <dbl> <chr>
1 accuracy    binary          0.579 Preprocessor1_Model1
2 roc_auc     binary          0.610 Preprocessor1_Model1
3 brier_class binary          0.241 Preprocessor1_Model1
```

##Here is the code for xgBoost:

```
# Recipe
gb_recipe <-
  recipe(
    los_long ~ .,
    data = mimic_other
  ) |>
  step_unknown(all_nominal(), -all_outcomes()) |> # Handle missing categorical values
  # mean imputation for numeric variable
  step_impute_mean(`heart rate`, `non invasive blood pressure systolic`, `non invasive blood
    `respiratory rate`, `temperature fahrenheit`,
    bicarbonate, chloride, creatinine, glucose, potassium,
    sodium, hematocrit, wbc) |>
  # mode imputation for catrgorical variable
  step_impute_mode(marital_status) |>
  # create traditional dummy variables
  step_dummy(all_nominal_predictors()) |>
  # zero-variance filter
  step_zv(all_numeric_predictors()) |>
  # center and scale numeric data
  step_normalize(all_numeric_predictors()) |>
  # estimate the means and standard deviations
  # prep(training = Heart_other, retain = TRUE) |>
  print()
```

-- Recipe -----

```
-- Inputs
```

```
Number of variables by role
```

```
outcome:    1  
predictor: 18
```

```
-- Operations
```

```
* Unknown factor level assignment for: all_nominal() -all_outcomes()
```

```
* Mean imputation for: `heart rate`, ...
```

```
* Mode imputation for: marital_status
```

```
* Dummy variables from: all_nominal_predictors()
```

```
* Zero variance filter on: all_numeric_predictors()
```

```
* Centering and scaling for: all_numeric_predictors()
```

```
# Model  
gb_mod <-  
  boost_tree(  
    mode = "classification",  
    trees = 1000,  
    tree_depth = tune(),  
    learn_rate = tune()  
  ) |>  
  set_engine("xgboost")  
gb_mod
```

Boosted Tree Model Specification (classification)

Main Arguments:

```
trees = 1000
tree_depth = tune()
learn_rate = tune()
```

Computational engine: xgboost

```
# bundle recipe and model(workflow step)
gb_wf <- workflow() |>
  add_recipe(gb_recipe) |>
  add_model(gb_mod)
gb_wf
```

== Workflow =====

Preprocessor: Recipe

Model: boost_tree()

-- Preprocessor -----

6 Recipe Steps

```
* step_unknown()
* step_impute_mean()
* step_impute_mode()
* step_dummy()
* step_zv()
* step_normalize()
```

-- Model -----

Boosted Tree Model Specification (classification)

Main Arguments:

```
trees = 1000
tree_depth = tune()
learn_rate = tune()
```

Computational engine: xgboost

```
# Tune
param_grid <- grid_regular(
  tree_depth(range = c(1L, 3L)),
```

```

  learn_rate(range = c(-5, 2), trans = log10_trans()),
  levels = c(3, 3)
)
param_grid

```

```

# A tibble: 9 x 2
  tree_depth learn_rate
    <int>      <dbl>
1         1  0.00001
2         2  0.00001
3         3  0.00001
4         1  0.0316
5         2  0.0316
6         3  0.0316
7         1 100
8         2 100
9         3 100

```

```

#CV
set.seed(203)

folds <- vfold_cv(mimic_other, v = 2)
folds

```

```

# 2-fold cross-validation
# A tibble: 2 x 2
  splits          id
  <list>        <chr>
1 <split [23610/23611]> Fold1
2 <split [23611/23610]> Fold2

```

```

gb_fit <- gb_wf |>
  tune_grid(
    resamples = folds,
    grid = param_grid,
    metrics = metric_set(roc_auc, accuracy)
  )
gb_fit

```

```

# Tuning results

```

```
# 2-fold cross-validation
```

```
# A tibble: 2 x 4
```

```
  splits          id    .metrics          .notes
  <list>         <chr> <list>         <list>
1 <split [23610/23611]> Fold1 <tibble [18 x 6]> <tibble [0 x 3]>
2 <split [23611/23610]> Fold2 <tibble [18 x 6]> <tibble [0 x 3]>
```

```
# Visualize
```

```
gb_fit |>
  collect_metrics() |>
  print(width = Inf) |>
  filter(.metric == "roc_auc") |>
  ggplot(mapping = aes(x = learn_rate, y = mean, color = factor(tree_depth))) +
  geom_point() +
  labs(x = "Learning Rate", y = "CV AUC") +
  scale_x_log10()
```

```
# A tibble: 18 x 8
```

	tree_depth	learn_rate	.metric	.estimator	mean	n	std_err
	<int>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>
1	1	0.00001	accuracy	binary	0.537	2	0.00482
2	1	0.00001	roc_auc	binary	0.537	2	0.00383
3	2	0.00001	accuracy	binary	0.545	2	0.00200
4	2	0.00001	roc_auc	binary	0.549	2	0.000879
5	3	0.00001	accuracy	binary	0.547	2	0.00440
6	3	0.00001	roc_auc	binary	0.562	2	0.000426
7	1	0.0316	accuracy	binary	0.575	2	0.00465
8	1	0.0316	roc_auc	binary	0.609	2	0.00469
9	2	0.0316	accuracy	binary	0.581	2	0.00459
10	2	0.0316	roc_auc	binary	0.617	2	0.00403
11	3	0.0316	accuracy	binary	0.585	2	0.00548
12	3	0.0316	roc_auc	binary	0.619	2	0.00413
13	1	100	accuracy	binary	0.537	2	0.00482
14	1	100	roc_auc	binary	0.537	2	0.00383
15	2	100	accuracy	binary	0.545	2	0.00200
16	2	100	roc_auc	binary	0.546	2	0.00421
17	3	100	accuracy	binary	0.500	2	0.0425
18	3	100	roc_auc	binary	0.505	2	0.0497

```
.config
```

```
<chr>
```

```
1 Preprocessor1_Model1
```


```
2 Preprocessor1_Model1
```



```

3 Preprocessor1_Model2
4 Preprocessor1_Model2
5 Preprocessor1_Model3
6 Preprocessor1_Model3
7 Preprocessor1_Model4
8 Preprocessor1_Model4
9 Preprocessor1_Model5
10 Preprocessor1_Model5
11 Preprocessor1_Model6
12 Preprocessor1_Model6
13 Preprocessor1_Model7
14 Preprocessor1_Model7
15 Preprocessor1_Model8
16 Preprocessor1_Model8
17 Preprocessor1_Model9
18 Preprocessor1_Model9

```



```

gb_fit |>
  show_best(metric = "roc_auc")

```

```

# A tibble: 5 x 8
  tree_depth learn_rate .metric .estimator mean      n std_err .config

```

	<int>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	3	0.0316	roc_auc	binary	0.619	2	0.00413	Preprocessor1_M~
2	2	0.0316	roc_auc	binary	0.617	2	0.00403	Preprocessor1_M~
3	1	0.0316	roc_auc	binary	0.609	2	0.00469	Preprocessor1_M~
4	3	0.00001	roc_auc	binary	0.562	2	0.000426	Preprocessor1_M~
5	2	0.00001	roc_auc	binary	0.549	2	0.000879	Preprocessor1_M~

```
best_gb <- gb_fit |>
  select_best(metric = "roc_auc")
best_gb
```

```
# A tibble: 1 x 3
  tree_depth learn_rate .config
    <int>      <dbl> <chr>
1         3      0.0316 Preprocessor1_Model6
```

```
#finalize model
final_wf <- gb_wf |>
  finalize_workflow(best_gb)
final_wf
```

```
== Workflow =====
Preprocessor: Recipe
Model: boost_tree()

-- Preprocessor -----
6 Recipe Steps

* step_unknown()
* step_impute_mean()
* step_impute_mode()
* step_dummy()
* step_zv()
* step_normalize()

-- Model -----
Boosted Tree Model Specification (classification)

Main Arguments:
  trees = 1000
  tree_depth = 3
```

```
learn_rate = 0.0316227766016838
```

Computational engine: xgboost

```
final_fit <-  
  final_wf |>  
  last_fit(data_split)  
final_fit
```

Resampling results

Manual resampling

A tibble: 1 x 6

	splits	id	.metrics	.notes	.predictions	.workflow
	<list>	<chr>	<list>	<list>	<list>	<list>
1	<split [47221/47223]>	train/test sp~	<tibble>	<tibble>	<tibble>	<workflow>

```
final_fit |> collect_metrics()
```

A tibble: 3 x 4

	.metric	.estimator	.estimate	.config
	<chr>	<chr>	<dbl>	<chr>
1	accuracy	binary	0.584	Preprocessor1_Model1
2	roc_auc	binary	0.622	Preprocessor1_Model1
3	brier_class	binary	0.238	Preprocessor1_Model1

##Model Stacking code:

```
# set up 2-fold CV  
set.seed(203)  
library(stacks)  
folds <- vfold_cv(mimic_other, v = 2)  
  
# Base models  
# logistic regression  
logit_mod <-  
  logistic_reg(  
    penalty = tune(),  
    mixture = tune()  
  ) |>  
  set_engine("glmnet", standardize = FALSE)  
logit_mod
```

Logistic Regression Model Specification (classification)

Main Arguments:

```
penalty = tune()  
mixture = tune()
```

Engine-Specific Arguments:

```
standardize = FALSE
```

Computational engine: glmnet

```
logit_wf <- workflow() |>  
  add_recipe(gb_recipe) |>  
  add_model(logit_mod)  
logit_wf
```

== Workflow =====

Preprocessor: Recipe

Model: logistic_reg()

-- Preprocessor -----

6 Recipe Steps

```
* step_unknown()  
* step_impute_mean()  
* step_impute_mode()  
* step_dummy()  
* step_zv()  
* step_normalize()
```

-- Model -----

Logistic Regression Model Specification (classification)

Main Arguments:

```
penalty = tune()  
mixture = tune()
```

Engine-Specific Arguments:

```
standardize = FALSE
```

Computational engine: glmnet

```
logit_grid <- grid_regular(
  penalty(range = c(-6, 3)),
  mixture(),
  levels = c(100, 5)
)

logit_res <- logit_wf |>
  tune_grid(
    resamples = folds,
    grid = logit_grid,
    control = control_stack_grid()
  )
```

i The workflow being saved contains a recipe, which is 5.85 Mb in i memory. If this was not intentional, please set the control setting i `save_workflow = FALSE`.

```
logit_res
```

```
# Tuning results
# 2-fold cross-validation
# A tibble: 2 x 5
  splits          id    .metrics          .notes          .predictions
  <list>         <chr> <list>         <list>         <list>
1 <split [23610/23611]> Fold1 <tibble [1,500 x 6]> <tibble [0 x 3]> <tibble>
2 <split [23611/23610]> Fold2 <tibble [1,500 x 6]> <tibble [0 x 3]> <tibble>
```

```
#random forest
library(ranger)
rf_mod <-
  rand_forest(
    mode = "classification",
    mtry = tune(),
    trees = tune()
  ) |>
  set_engine("ranger",
    importance = "impurity")
rf_mod
```

Random Forest Model Specification (classification)

Main Arguments:

```
mtry = tune()
trees = tune()
```

Engine-Specific Arguments:

```
importance = impurity
```

Computational engine: ranger

```
#i used xgboost recipe here since they are essentially the same
rf_wf <- workflow() |>
  add_recipe(gb_recipe) |>
  add_model(rf_mod)
rf_wf
```

== Workflow =====

Preprocessor: Recipe

Model: rand_forest()

-- Preprocessor -----

6 Recipe Steps

```
* step_unknown()
* step_impute_mean()
* step_impute_mode()
* step_dummy()
* step_zv()
* step_normalize()
```

-- Model -----

Random Forest Model Specification (classification)

Main Arguments:

```
mtry = tune()
trees = tune()
```

Engine-Specific Arguments:

```
importance = impurity
```

Computational engine: ranger

```
rf_grid <- grid_regular(
  trees(range = c(100L, 500L)),
  mtry(range = c(1L, 5L)),
  levels = c(5, 5)
)

rf_res <- tune_grid(
  object = rf_wf,
  resamples = folds,
  grid = rf_grid,
  control = control_stack_grid()
)
```

i The workflow being saved contains a recipe, which is 5.85 Mb in i memory. If this was not intentional, please set the control setting i `save_workflow = FALSE`.

```
rf_res
```

```
# Tuning results
# 2-fold cross-validation
# A tibble: 2 x 5
  splits          id    .metrics          .notes          .predictions
  <list>         <chr> <list>          <list>          <list>
1 <split [23610/23611]> Fold1 <tibble [75 x 6]> <tibble [0 x 3]> <tibble>
2 <split [23611/23610]> Fold2 <tibble [75 x 6]> <tibble [0 x 3]> <tibble>
```

```
#xgboost
gb_mod <- boost_tree(
  mode = "classification",
  trees = 1000,
  tree_depth = tune(),
  learn_rate = tune()
) |> set_engine("xgboost")
gb_mod
```

Boosted Tree Model Specification (classification)

Main Arguments:
trees = 1000

```
tree_depth = tune()
learn_rate = tune()
```

Computational engine: xgboost

```
gb_wf <- workflow() |>
  add_recipe(gb_recipe) |>
  add_model(gb_mod)
gb_wf
```

== Workflow =====

Preprocessor: Recipe

Model: boost_tree()

-- Preprocessor -----

6 Recipe Steps

```
* step_unknown()
* step_impute_mean()
* step_impute_mode()
* step_dummy()
* step_zv()
* step_normalize()
```

-- Model -----

Boosted Tree Model Specification (classification)

Main Arguments:

```
trees = 1000
tree_depth = tune()
learn_rate = tune()
```

Computational engine: xgboost

```
gb_grid <- grid_regular(
  tree_depth(range = c(1L, 3L)),
  learn_rate(range = c(-5, 2), trans = log10_trans()),
  levels = c(3, 3)
)
gb_grid
```

A tibble: 9 x 2

	tree_depth	learn_rate
	<int>	<dbl>
1	1	0.00001
2	2	0.00001
3	3	0.00001
4	1	0.0316
5	2	0.0316
6	3	0.0316
7	1	100
8	2	100
9	3	100

```
gb_res <-
  tune_grid(
    object = gb_wf,
    resamples = folds,
    grid = gb_grid,
    control = control_stack_grid()
  )
```

i The workflow being saved contains a recipe, which is 5.85 Mb in i memory. If this was not intentional, please set the control setting i `save_workflow = FALSE`.

```
gb_res
```

```
# Tuning results
# 2-fold cross-validation
# A tibble: 2 x 5
  splits          id  .metrics          .notes          .predictions
  <list>         <chr> <list>          <list>          <list>
1 <split [23610/23611]> Fold1 <tibble [27 x 6]> <tibble [0 x 3]> <tibble>
2 <split [23611/23610]> Fold2 <tibble [27 x 6]> <tibble [0 x 3]> <tibble>
```

```
#stacking models
mimic_model_st <- stacks() |>
  add_candidates(logit_res)|>
  add_candidates(rf_res) |>
  add_candidates(gb_res) |>
  blend_predictions(
    penalty = 10^(-6:2),
```

```

    metrics = c("roc_auc")
  ) |>
  fit_members()

```

Warning: Predictions from 742 candidates were identical to those from existing candidates and were removed from the data stack.

Warning: The `...` are not used in this function but one or more arguments were passed: 'metrics'

```
mimic_model_st
```

```
-- A stacked ensemble model -----
```

Out of 163 possible candidate members, the ensemble retained 15.

Penalty: 0.001.

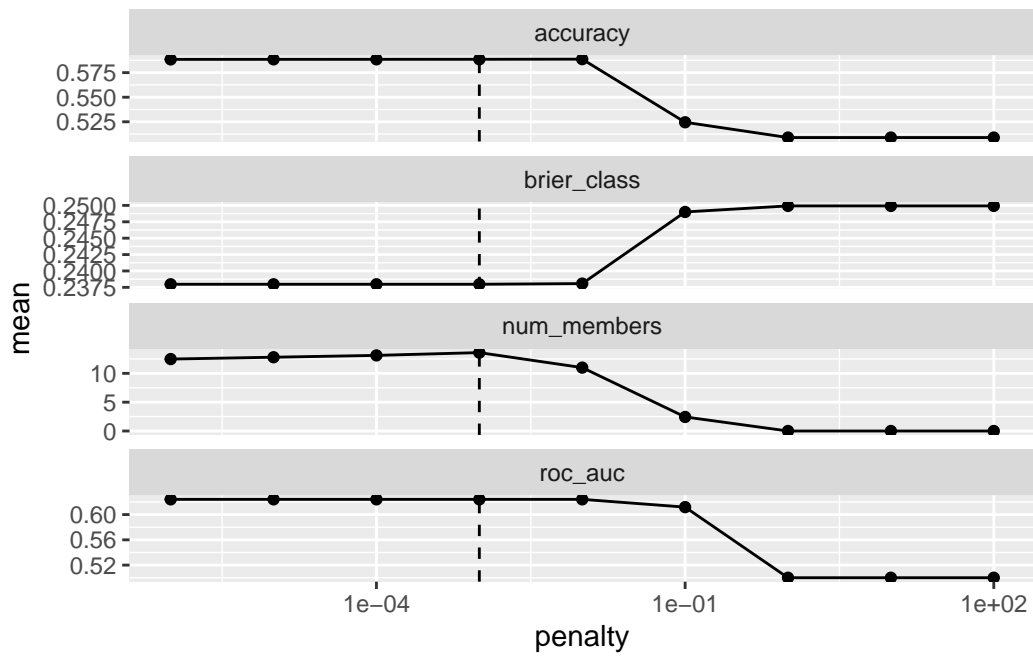
Mixture: 1.

The 10 highest weighted member classes are:


```
# A tibble: 10 x 3
```

	member <chr>	type <chr>	weight <dbl>
1	.pred_yes_gb_res_1_6	boost_tree	1.83
2	.pred_yes_rf_res_1_08	rand_forest	0.825
3	.pred_yes_rf_res_1_20	rand_forest	0.601
4	.pred_yes_rf_res_1_17	rand_forest	0.452
5	.pred_yes_rf_res_1_07	rand_forest	0.408
6	.pred_yes_rf_res_1_21	rand_forest	0.348
7	.pred_yes_rf_res_1_06	rand_forest	0.223
8	.pred_yes_logit_res_1_101	logistic_reg	0.118
9	.pred_yes_logit_res_1_201	logistic_reg	0.0928
10	.pred_yes_rf_res_1_09	rand_forest	0.0871

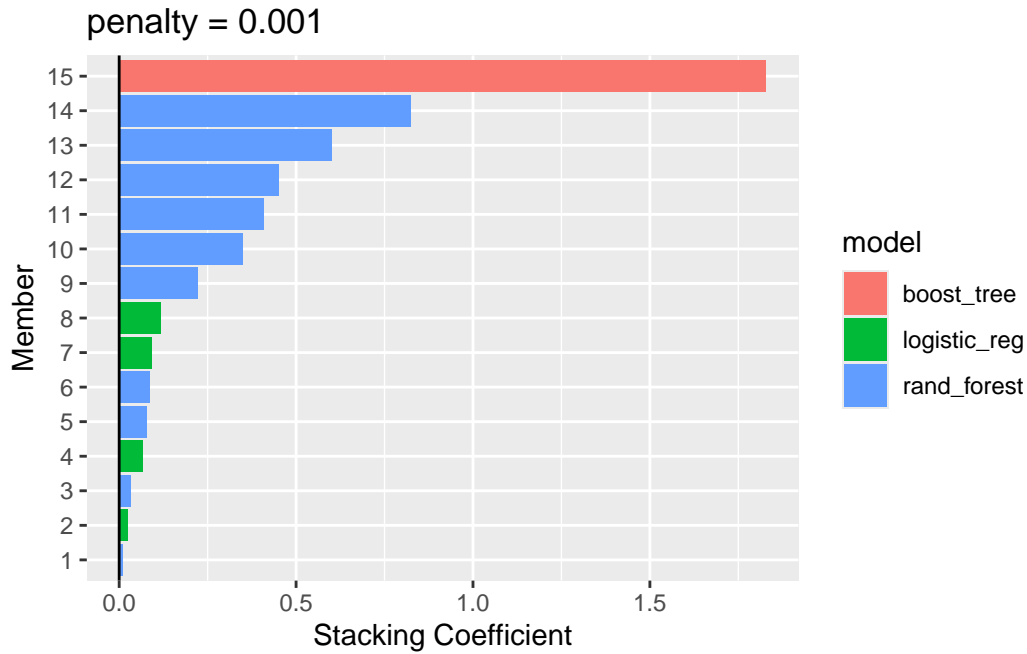
```
autoplot(mimic_model_st)
```



```
autoplot(mimic_model_st, type = "members")
```



```
autoplot(mimic_model_st, type = "weights")
```



```
#collect_metrics(mimic_model_st, "rf_res")
```

```
collect_parameters(mimic_model_st, "rf_res")
```

```
# A tibble: 25 x 5
  member      mtry trees terms      coef
  <chr>      <int> <int> <chr>    <dbl>
1 rf_res_1_01      1   100 .pred_yes_rf_res_1_01 0
2 rf_res_1_02      1   200 .pred_yes_rf_res_1_02 0
3 rf_res_1_03      1   300 .pred_yes_rf_res_1_03 0
4 rf_res_1_04      1   400 .pred_yes_rf_res_1_04 0
5 rf_res_1_05      1   500 .pred_yes_rf_res_1_05 0
6 rf_res_1_06      2   100 .pred_yes_rf_res_1_06 0.223
7 rf_res_1_07      2   200 .pred_yes_rf_res_1_07 0.408
8 rf_res_1_08      2   300 .pred_yes_rf_res_1_08 0.825
9 rf_res_1_09      2   400 .pred_yes_rf_res_1_09 0.0871
10 rf_res_1_10     2   500 .pred_yes_rf_res_1_10 0
# i 15 more rows
```

```
#final classification
mimic_pred <- mimic_test %>%
  bind_cols(predict(mimic_model_st, ., type = "prob")) %>%
  print(width = Inf)
```

```
# A tibble: 47,223 x 21
```

	first_careunit	gender	age_at_intime	
	<fct>	<fct>	<int>	
1	Medical Intensive Care Unit (MICU)	F	52	
2	Medical/Surgical Intensive Care Unit (MICU/SICU)	F	46	
3	Cardiac Vascular Intensive Care Unit (CVICU)	F	57	
4	Other	M	56	
5	Medical Intensive Care Unit (MICU)	F	83	
6	Medical/Surgical Intensive Care Unit (MICU/SICU)	F	82	
7	Medical Intensive Care Unit (MICU)	F	81	
8	Other	M	90	
9	Other	M	53	
10	Cardiac Vascular Intensive Care Unit (CVICU)	F	58	
	marital_status	race	`heart rate`	`non invasive blood pressure systolic`
	<fct>	<fct>	<dbl>	<dbl>
1	WIDOWED	WHITE	91	84
2	MARRIED	WHITE	86	73
3	SINGLE	Other	80	104
4	<NA>	Other	111	112
5	MARRIED	WHITE	71	126
6	MARRIED	WHITE	71	126
7	WIDOWED	WHITE	124	87
8	WIDOWED	WHITE	96	108
9	SINGLE	WHITE	106	140
10	<NA>	WHITE	80	109
			`non invasive blood pressure diastolic`	`respiratory rate`
			<dbl>	<dbl>
1			48	24
2			56	19
3			70	14
4			80	22
5			61	18
6			61	18
7			42	25
8			61	26
9			99	12
10			72	17

	`temperature fahrenheit`	bicarbonate	chloride	creatinine	glucose	potassium
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	98.7	25	95	0.7	102	6.7
2	97.7	NA	98	NA	NA	4.1
3	97.2	24	102	0.9	288	3.5
4	97.9	18	NA	3.1	95	6.5
5	95.9	26	85	1.4	133	5.7
6	95.9	23	98	2.8	117	4.9
7	103.	27	111	0.6	173	4.4
8	98.1	23	102	1.9	105	4.4
9	96.7	18	106	0.9	269	5.3
10	99	NA	NA	NA	NA	NA

	sodium	hematocrit	wbc	los_long	.pred_no	.pred_yes
	<dbl>	<dbl>	<dbl>	<fct>	<dbl>	<dbl>
1	126	41.1	6.9	no	0.520	0.480
2	139	NA	NA	no	0.492	0.508
3	137	34.9	7.2	no	0.614	0.386
4	125	34.3	16.8	yes	0.359	0.641
5	120	22.4	9.8	no	0.548	0.452
6	135	25.5	17.9	yes	0.556	0.444
7	144	34.7	10.5	yes	0.382	0.618
8	140	29.9	5.1	yes	0.473	0.527
9	135	43.1	16.9	yes	0.485	0.515
10	NA	NA	NA	no	0.323	0.677

i 47,213 more rows

```
yardstick::roc_auc(
  mimic_pred,
  truth = los_long,
  contains(".pred_No")
)
```

```
# A tibble: 1 x 3
  .metric .estimator .estimate
  <chr>   <chr>         <dbl>
1 roc_auc binary       0.629
```

```
mimic_pred <- mimic_test |>
  select(los_long) |>
  bind_cols(
    predict(mimic_model_st,
            mimic_test,
```

```

      type = "class",
      members = TRUE)) |>
print(width = Inf)

```

A tibble: 47,223 x 17

	los_long	.pred_class	.pred_class_logit_res_1_001	.pred_class_logit_res_1_101
	<fct>	<fct>	<fct>	<fct>
1	no	no	no	no
2	no	yes	no	no
3	no	no	no	no
4	yes	yes	yes	yes
5	no	no	no	no
6	yes	no	no	no
7	yes	yes	yes	yes
8	yes	yes	yes	yes
9	yes	yes	no	no
10	no	yes	yes	yes
	.pred_class_logit_res_1_201	.pred_class_logit_res_1_301		
	<fct>	<fct>		
1	no	no		
2	no	no		
3	no	no		
4	yes	yes		
5	no	no		
6	no	no		
7	yes	yes		
8	yes	yes		
9	no	no		
10	yes	yes		
	.pred_class_rf_res_1_06	.pred_class_rf_res_1_07	.pred_class_rf_res_1_08	
	<fct>	<fct>	<fct>	
1	no	no	no	
2	yes	yes	yes	
3	no	no	no	
4	yes	yes	yes	
5	no	no	no	
6	no	no	no	
7	yes	yes	yes	
8	no	yes	no	
9	no	no	no	
10	yes	yes	yes	
	.pred_class_rf_res_1_09	.pred_class_rf_res_1_13	.pred_class_rf_res_1_17	

	<fct>	<fct>	<fct>
1	yes	yes	yes
2	yes	yes	no
3	no	no	no
4	yes	yes	yes
5	yes	yes	no
6	no	yes	no
7	yes	yes	yes
8	yes	no	no
9	no	no	yes
10	yes	yes	yes
	.pred_class_rf_res_1_19	.pred_class_rf_res_1_20	.pred_class_rf_res_1_21
	<fct>	<fct>	<fct>
1	yes	yes	no
2	yes	yes	no
3	no	no	no
4	yes	yes	yes
5	no	no	yes
6	no	yes	yes
7	yes	yes	yes
8	yes	no	no
9	no	no	yes
10	yes	yes	yes
	.pred_class_rf_res_1_22	.pred_class_gb_res_1_6	
	<fct>	<fct>	
1	yes	no	
2	no	yes	
3	no	no	
4	yes	yes	
5	no	no	
6	no	no	
7	yes	yes	
8	no	yes	
9	no	yes	
10	yes	yes	

i 47,213 more rows

```
map(colnames(mimic_pred),
  ~mean(mimic_pred$los_long == pull(mimic_pred, .x))
) |>
set_names(colnames(mimic_pred)) |>
as_tibble() |>
```



```
pivot_longer(c(everything()), -los_long))
```

```
# A tibble: 16 x 3
  los_long name          value
  <dbl> <chr>          <dbl>
1      1 .pred_class      0.591
2      1 .pred_class_logit_res_1_001 0.561
3      1 .pred_class_logit_res_1_101 0.561
4      1 .pred_class_logit_res_1_201 0.561
5      1 .pred_class_logit_res_1_301 0.561
6      1 .pred_class_rf_res_1_06     0.581
7      1 .pred_class_rf_res_1_07     0.584
8      1 .pred_class_rf_res_1_08     0.584
9      1 .pred_class_rf_res_1_09     0.587
10     1 .pred_class_rf_res_1_13     0.588
11     1 .pred_class_rf_res_1_17     0.587
12     1 .pred_class_rf_res_1_19     0.587
13     1 .pred_class_rf_res_1_20     0.587
14     1 .pred_class_rf_res_1_21     0.581
15     1 .pred_class_rf_res_1_22     0.584
16     1 .pred_class_gb_res_1_6      0.584
```

4. Compare model classification performance on the test set. Report both the area under ROC curve and accuracy for each machine learning algorithm and the model stacking. Interpret the results. What are the most important features in predicting long ICU stays? How do the models compare in terms of performance and interpretability?

Solution:

```
library(vip)
```

Attaching package: 'vip'

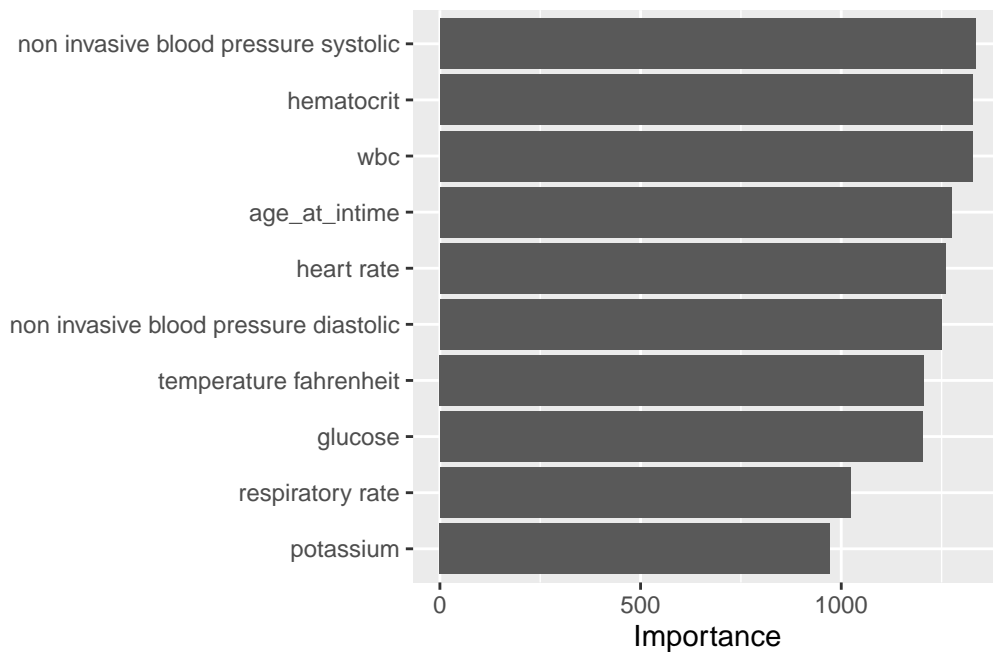
The following object is masked from 'package:utils':

vi

```
best_rf <- rf_res |> select_best(metric = "roc_auc")

rf_fit <- rf_wf |>
  finalize_workflow(best_rf) |>
  fit(data = mimic_other)

rf_fit |>
  extract_fit_parsnip() |>
  vip(num_features = 10)
```



Logistic regression gives an AUC of 0.586;

SVM gives an AUC of 0.596;

xgboost gives the best performance with an AUC of 0.622.

model stacking(logistic + random forest + xgboost) gives an AUC of 0.591. boosting is the dominant model among all three.

xgboost provides the best result so far. The most important features predicting `los_long` according to rf(since it has the highest weight in stacking) are: hematocrit, wbc, glucose, blood pressures, `age_at_intime` and heart rate.

xgboost performed best in ROC-AUC, meaning it distinguishes long ICU stays most effectively.

Logistic regression provides the best interpretability since variable coefficients indicates the feature impact directly.